

STATISTICAL DATA ANALYSIS and MACHINE LEARNING

July 14, 2020

1 INTRODUCTION AND SUMMARY

Machine learning models are based on statistics, optimization and computer science. Our project aims to compare the performance of a few machine learning models on different databases. For this analysis, we used the Python programming language on Jupyter notebooks. The project will be realized in 2 terms and we conducted the experimental part in our first period. We will present the theoretical discussion of mathematical basis of the models use in the fall period of 2021.

1.1 PROJECT PLAN

1.1.1 Aim of the Project

We aim to perform rigorous statistical analysis of 6 different machine learning models on 3 different datasets.

1.1.2 Scope of the Project

Within the scope of this project, 6 different machine learning models were selected from well-known and commonly used python libraries. A deeper theoretical exploration about the models will be given at the thesis stage. Since we have devoted our first period to experiments, we will show an explanatory code section of the project in this report.

1.1.3 Area of Usage

Within the scope of the analysis, the final version of the project can be used as a resource for those who want to learn the mathematical background of machine learning models, and how they can be applied to real-world data sets. We aim that the thesis will be a resource for those who want to learn from the beginning which models work best for which dataset.

1.1.4 Schedule

Task	Alloted Time
Research and discovery of databases	2 weeks
Clearing and preparing datasets	2 weeks
Determining the models and libraries to be used	1 week
Coding phase and evaluation of outputs	10 weeks
General arrangement of the code stage and annotation	2 weeks
Writing the report	3 weeks

1.1.5 Resources

All databases, libraries and models used are open source. All relevant references will be made to the books, articles and theses used for theory in the reference section in the thesis section. All work is stored on github. (<https://github.com/Code-Cash/ahmet>)

2 EXPERIMENTS

Our project is based on a thorough and rigorous statistical analysis of 6 different machine learning algorithms over 3 databases.

The code block you see below is the experimental stage of this thesis. We give descriptions about each cell just above the corresponding cell.

2.1 Loading the libraries

At the beginning of our project, we need to add the libraries required for our project.

```
[2]: import pandas as pd
import numpy as np
from pandas import Series, DataFrame
from sklearn.preprocessing import scale
from sklearn.model_selection import cross_val_score
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from PIL import Image
from sklearn.metrics import accuracy_score

import statsmodels.api as sm
from statsmodels.formula.api import ols
from statsmodels.stats.anova import AnovaRM

from matplotlib import pyplot as plt
from matplotlib import pyplot
import seaborn as sns
import xgboost as xgb
from sklearn.naive_bayes import GaussianNB

from xgboost import XGBClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

from sklearn.feature_selection import RFE
```

2.2 Loading the models

We load 6 blank machine learning models after adding our libraries.

```
[3]: dt = DecisionTreeClassifier()
      rf = RandomForestClassifier(n_estimators=100, random_state=0, n_jobs = -1)
      logreg = LogisticRegression(solver='liblinear',multi_class='ovr')
      svm = SVC(gamma='auto')
      xgb = xgb.XGBClassifier()
      gnb = GaussianNB()
```

2.3 Loading the first data set

We add our first data set using the pandas library. Then we display the first 5 elements of our dataset using the panda's head function.

```
[5]: database1 = pd.read_csv("data.csv", sep=';')
      database1.head()
```

```
[5]:   MUSK   2   3   4   5   6   7   8   9  10  ...  159  160  161  162  163  \
0  MUSK  46 -108 -60 -69 -117  49  38 -161 -8  ... -308   52   -7   39  126
1  MUSK  41 -188 -145  22 -117  -6  57 -171 -39  ...  -59   -2   52  103  136
2  MUSK  46 -194 -145  28 -117  73  57 -168 -39  ... -134 -154   57  143  142
3  MUSK  41 -188 -145  22 -117  -7  57 -170 -39  ...  -60   -4   52  104  136
4  MUSK  41 -188 -145  22 -117  -7  57 -170 -39  ...  -60   -4   52  104  137

      164  165  166  167  168
0  156  -50 -112   96  1.0
1  169  -61 -136   79  1.0
2  165  -67 -145   39  1.0
3  168  -60 -135   80  1.0
4  168  -60 -135   80  1.0

[5 rows x 168 columns]
```

2.4 Feature selection and splitting the train and test data sets

Let us denote the features we are going to use with **X** and response variable with **y**. With the drop function, we get the columns other than the columns 'musk' and '168'. For the response variable, we just take the column named '168'. Then we determine the size of our training and test sets. For this dataset, 90% of all data is used for training and the remaining 10% is used for testing using the function "train_test_split".

```
[7]: X = database1.drop(["MUSK", "168"],axis=1)
      y = database1['168']
      X_train, X_test, y_train, y_test = train_test_split(X, y,test_size=0.10 )
```

2.5 A generic function for training and evaluating a model

At this stage, we create a function that will run our 6 algorithms for our first dataset and give us 3 outputs.

If we go into detail: first, we train our model with the features **X**, which is the training group we mentioned in Section 4, with the features in **y** as dependent variable. Then we make a prediction with the predict function according to the test data set. We record this prediction as **y_pred**. The test data set **y_test** is the test set of the results will give us the accuracy of the model. Then we print the classification report and confusion matrix with the help of the scikit-learn library. The explanation of all these functions will be found in the theory section of our thesis.

```
[4]: def model(name, X_train, y_train, X_test, y_test):  
    name.fit(X_train, y_train)  
    y_pred = name.predict(X_test)  
    score= accuracy_score(y_test, y_pred) * 100  
    print(str(name) + "Accuracy:",score)  
    report_name=classification_report(y_test, y_pred)  
    print(report_name)  
    print(confusion_matrix(y_test,y_pred))
```

2.6 Recursive Feature Elimination

The purpose of the function below is to find the most suitable features for us in our database. The **rfe** (Recursive Feature Elimination) function produces an output based on the model we want and the number of desired features. It shows them in order, according to their effect.

```
[5]: def rfe(model,X,y):  
    rfe = RFE(model, 16)  
    fit = rfe.fit(X, y)  
    print("Num Features: %s" % (fit.n_features_))  
    print("Selected Features: %s" % (fit.support_))  
    print("Feature Ranking: %s" % (fit.ranking_))
```

2.7 Cross validation

Cross validation calculates the average accuracy and standard deviation for each part by dividing our data into as many parts as we want. This gives us an important idea of how well the model works.

```
[6]: def crossval(model,X,y,n=10):  
    scores=cross_val_score(model, X, y, cv=n, scoring ="accuracy")  
    print(str(model) + "Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), scores.  
→std() * 2))
```

2.8 Experiments on the first data set

2.8.1 Decision Tree

Let us evaluate the Decision Tree model on our first data set.

```
[15]: model(dt,X_train,y_train,X_test,y_test)
      crossval(dt,X,y,n)
      rfe(dt,X,y)
```

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                      max_depth=None, max_features=None, max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, presort='deprecated',
                      random_state=None, splitter='best')Accuracy:
```

```
96.21212121212122
```

	precision	recall	f1-score	support
0.0	0.97	0.98	0.98	549
1.0	0.90	0.87	0.89	111
accuracy			0.96	660
macro avg	0.94	0.93	0.93	660
weighted avg	0.96	0.96	0.96	660

```
[[538 11]
 [ 14 97]]
```

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                      max_depth=None, max_features=None, max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, presort='deprecated',
                      random_state=None, splitter='best')Accuracy: 0.78 (+/-
```

```
0.40)
```

```
Num Features: 16
```

```
Selected Features: [False False False False False False False False  True False
False False
```

```
False False False False False False False False False False False False
False False False False False False False  True False False False  True
False  True False False False False  True False False False False False
False  True False False False False False False False False False False
False False False False  True False False False False False  True False
False False False False False False False False False False False False
False False False False False False False False False False False False
False False False False False False False False False False False False
False  True False  True False  True False False False False False  True
False False False False False False False  True False False False False
False False False False False False  True False False False False False
False False False False False False  True False False False]
```

```
Feature Ranking: [146 145 45 62 46 32 66 61 1 53 71 70 75 60 78
40 9 69
```

```

83 15 93 12 90 88 24 103 8 108 43 110 85 1 122 41 34 1
54 1 33 74 80 18 1 136 141 86 134 36 56 1 139 3 68 2
6 37 95 98 92 58 38 82 21 26 140 1 117 50 35 118 128 100
144 130 17 143 1 113 119 123 121 149 1 142 4 138 89 31 5 137
77 106 14 59 84 16 91 76 11 73 25 65 27 63 109 10 81 28
64 107 51 67 120 105 101 30 22 96 114 52 29 1 97 1 99 1
127 147 129 47 49 1 7 125 124 44 102 94 111 1 19 48 39 135
132 133 87 126 115 131 1 23 72 116 104 42 112 13 79 55 57 20
1 148 150 151]

```

2.8.2 Random Forest

Let us evaluate the Random Forest algorithm on our first data set.

```
[17]: model(rf,X_train,y_train,X_test,y_test)
      crossval(rf,X,y)
      rfe(rf,X,y)
```

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=None, max_features='auto',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=100,
                        n_jobs=-1, oob_score=False, random_state=0, verbose=0,
                        warm_start=False)Accuracy: 97.12121212121212
```

	precision	recall	f1-score	support
0.0	0.97	1.00	0.98	549
1.0	0.99	0.84	0.91	111
accuracy			0.97	660
macro avg	0.98	0.92	0.95	660
weighted avg	0.97	0.97	0.97	660

```
[[548 1]
 [ 18 93]]
```

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=None, max_features='auto',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=100,
                        n_jobs=-1, oob_score=False, random_state=0, verbose=0,
                        warm_start=False)Accuracy: 0.80 (+/- 0.41)
```

Num Features: 16

Selected Features: [False False False False False False False False True False False False]

```

False False False False False False False False False False False False
False False False False False False False False False False False True
False False False False False False False False False False False False
False True False True False False False False False False False False
False False False False False True False False False False False False
False False False False False False False False False False False False
False False False False False False False True False False True False
False False False False False False False False False False False False
False True False False False False False False False False False False
False True False True False True False False False False False True
False False False False False False False True False False False False
False False False False False False True False False False False False
False False False False False True True False False False]
Feature Ranking: [ 7 143 50 72 150 77 87 83 1 20 110 136 45 19 128
145 28 123
38 44 2 8 58 121 17 95 131 134 94 96 106 25 43 53 24 1
140 40 74 130 76 29 5 69 4 125 144 68 86 1 49 1 56 98
32 59 147 90 42 105 10 97 9 92 116 1 151 47 138 113 108 31
122 126 62 149 91 82 23 39 89 22 12 75 11 16 139 119 26 114
127 1 61 81 1 27 120 60 71 141 115 13 55 63 35 93 73 46
15 1 66 30 102 99 37 36 65 54 135 112 84 1 52 1 41 1
103 57 78 64 51 1 18 80 67 129 137 101 104 1 21 132 70 33
142 88 148 118 117 79 1 124 85 100 107 111 146 109 133 3 48 1
1 34 6 14]

```

2.8.3 Logistic Regression

Let us evaluate the Logistic Regression algorithm on our first data set.

```
[18]: model(logreg,X_train,y_train,X_test,y_test)
      crossval(logreg,X,y)
      rfe(logreg,X,y)
```

```

LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='ovr', n_jobs=None, penalty='l2',
                    random_state=None, solver='liblinear', tol=0.0001, verbose=0,
                    warm_start=False)Accuracy: 95.0

```

	precision	recall	f1-score	support
0.0	0.95	0.99	0.97	549
1.0	0.92	0.77	0.84	111
accuracy			0.95	660
macro avg	0.94	0.88	0.90	660
weighted avg	0.95	0.95	0.95	660

```
[[542 7]
```

```
[ 26  85]]
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='ovr', n_jobs=None, penalty='l2',
                    random_state=None, solver='liblinear', tol=0.0001, verbose=0,
                    warm_start=False)Accuracy: 0.81 (+/- 0.28)

Num Features: 16
Selected Features: [False False False False False False False False False False
False False
                    True False False False False False False False False
False False True False False False False False False False False False
False False False False True False True False False False False False
False False False False False False False False True True False False
False False False False False False True False False False False True
False False False False False False False False False False False False
False False True False False False False False False False False False
False False True False False False False False False False False False
False False False False False False False False False False True False
False False False False False False False False False False False False
False False False False True False False False False False False False
False False False True False True True False False False]

Feature Ranking: [ 16  48  14  73 145  60 146  95  13  10  71  35  1  62  63
65  4 122
142 140 138  12 128 137  67  88  1 101  55 125  49  43  17 130  25  39
100  19  38  31  1  41  1 143  85 147  46  78  94  21  80  6  76 149
148  91  1  1  57 106 117 123 114 108  27  18  1 102 141  32  15  1
 70  84  87  82 124  64  93 139  75 121  81 116  9  45  1  92  97  59
134  51  54  69 151  2 133  5  1 103  40  52  74  86  61 127  83 104
111 112 110 129  77  29 105 107  72  90  1  11  66  56  58  98  3  50
 8 126  7  26  28 119  79  36  96  20  47  1  33  34 109 150  68  23
144 120  37 118  1  30  24  53 136  99  42  44 135 113 131  1 132  1
 1  22  89 115]
```

2.8.4 Support Vector Machines (SVM)

Let us evaluate the Support Vector Machines algorithm on our first data set.

```
[19]: model(svm,X_train,y_train,X_test,y_test)
      crossval(svm,X,y)
```

```
SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)Accuracy: 88.63636363636364
      precision    recall  f1-score   support

      0.0         0.88         1.00         0.94         549
```


	1.0	1.00	0.32	0.49	111
accuracy				0.89	660
macro avg		0.94	0.66	0.71	660
weighted avg		0.90	0.89	0.86	660

```
[[549  0]
 [ 75 36]]
```

```
SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)Accuracy: 0.85 (+/- 0.01)
```

2.8.5 XGBoost

Let us evaluate the XGBoost algorithm on our first data set.

```
[20]: model(xgb,X_train,y_train,X_test,y_test)
      crossval(xgb,X,y)
      rfe(xgb,X,y)
```

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
              importance_type='gain', interaction_constraints='',
              learning_rate=0.300000012, max_delta_step=0, max_depth=6,
              min_child_weight=1, missing=nan, monotone_constraints='()',
              n_estimators=100, n_jobs=0, num_parallel_tree=1,
              objective='binary:logistic', random_state=0, reg_alpha=0,
              reg_lambda=1, scale_pos_weight=1, subsample=1,
              tree_method='exact', validate_parameters=1,
              verbosity=None)Accuracy: 99.0909090909091
              precision    recall  f1-score   support
```

	0.0	0.99	1.00	0.99	549
	1.0	0.99	0.95	0.97	111
accuracy				0.99	660
macro avg		0.99	0.98	0.98	660
weighted avg		0.99	0.99	0.99	660

```
[[548  1]
 [  5 106]]
```

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
              importance_type='gain', interaction_constraints='',
              learning_rate=0.300000012, max_delta_step=0, max_depth=6,
              min_child_weight=1, missing=nan, monotone_constraints='()',
              n_estimators=100, n_jobs=0, num_parallel_tree=1,
```

```

        objective='binary:logistic', random_state=0, reg_alpha=0,
        reg_lambda=1, scale_pos_weight=1, subsample=1,
        tree_method='exact', validate_parameters=1,
verbosity=None)Accuracy: 0.86 (+/- 0.28)
Num Features: 16
Selected Features: [ True False False False False False False False  True False
False False
False False False False False False False False False  True False False
False False False False False False False  True False False  True  True
False False False False False False False False False False False False
False False False False False False False False False False False False
False False  True False False  True False False False False False False
False False False False False False False False False False False False
False False False False False False False  True False False False False
False False False False False False False False False False False False
False  True False False False False False False False False False False
False False False  True False  True False False False False False  True
False False False False False False False False False False False False
False False False False False False  True False False False False False
False False False False False False  True False  True False]
Feature Ranking: [ 1 47 51 66 50 49 92 125 1 3 116 113 102 99 87
95 6 70
80 121 10 1 75 109 53 150 31 81 77 63 18 1 67 40 1 1
111 4 119 38 148 15 29 74 127 136 128 82 94 16 62 21 84 73
8 46 103 48 36 142 34 11 1 132 135 1 27 39 90 93 147 122
120 104 97 105 9 139 133 20 85 112 26 98 144 43 151 35 23 118
56 1 19 107 114 54 134 88 78 124 24 37 57 89 65 79 106 110
55 1 32 108 117 33 52 41 44 71 91 72 64 14 58 1 83 1
140 126 129 69 5 1 17 123 115 42 28 61 130 2 59 145 96 45
25 13 146 60 143 149 1 141 12 76 30 7 138 101 100 86 131 22
1 68 1 137]

```

2.8.6 Gaussian Naive Bayes

Let us evaluate the Gaussian Naive Bayes algorithm on our first data set.

```
[13]: model(gnb,X_train,y_train,X_test,y_test)
      crossval(gnb,X,y)
```

```
GaussianNB(priors=None, var_smoothing=1e-09)Accuracy: 83.03030303030303
```

	precision	recall	f1-score	support
0.0	0.95	0.84	0.89	560
1.0	0.46	0.76	0.58	100
accuracy			0.83	660
macro avg	0.71	0.80	0.73	660
weighted avg	0.88	0.83	0.85	660

```
[[472  88]
 [ 24  76]]
GaussianNB(priors=None, var_smoothing=1e-09)Accuracy: 0.80 (+/- 0.31)
```

2.9 Experiments on Our Second Data Set

Our second dataset is slightly different from the first one. The training and test sets consist of two different csv blocks, so we define these two datasets separately.

```
[8]: database2 = pd.read_csv("shuttle-train.csv", sep=';')
      database2.head()
```

```
[8]:
```

	f1	f2	f3	f4	f5	f6	f7	f8	f9	d
0	50	21	77	0	28	0	27	48	22	2
1	55	0	92	0	0	26	36	92	56	4
2	53	0	82	0	52	-5	29	30	2	1
3	37	0	76	0	28	18	40	48	8	1
4	37	0	79	0	34	-26	43	46	2	1

```
[9]: database3 = pd.read_csv("shuttle-test.csv", sep=';')
      database3.head()
```

```
[9]:
```

	f1	f2	f3	f4	f5	f6	f7	f8	f9	d
0	55	0	81	0	-6	11	25	88	64	4
1	56	0	96	0	52	-4	40	44	4	4
2	50	-1	89	-7	50	0	39	40	2	1
3	53	9	79	0	42	-2	25	37	12	4
4	55	2	82	0	54	-6	26	28	2	1

2.9.1 Train and Test Split

At this stage, we already have separate test and train subsets, and as a result, we do not need to use the `train_test_split` function we used for the first dataset. Because currently our training and test sets are already given. All we have to do is to assign the independent and dependent variables.

```
[10]: X1_train = database2.drop(["d"], axis=1)
      X1_test = database3.drop(["d"], axis=1)
      y1_train = database2['d']
      y1_test = database3['d']
```

2.9.2 Combining Train and Test for Cross-Validation

This stage is actually combining the operations we did in stage 9, the two datasets above and then assigning the properties as `df_row_reindex_X` and the results as `df_row_reindex_y`. As a result, it allows us to work with a single dataset, not two. We do this with the function named `concat` in the pandas library.

```
[11]: df_row_reindex = pd.concat([database2, database3], ignore_index=True)
df_row_reindex_X = df_row_reindex.drop(["d"], axis=1)
df_row_reindex_y = df_row_reindex['d']
```

2.9.3 Decision Trees

Let us evaluate the Decision Tree algorithm on our second data set.

```
[13]: model(dt,X1_train,y1_train,X1_test,y1_test)
crossval(dt,df_row_reindex_X,df_row_reindex_y)
rfe(dt,df_row_reindex_X,df_row_reindex_y)
```

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                        max_depth=None, max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, presort='deprecated',
                        random_state=None, splitter='best')Accuracy:
```

99.99310344827586

	precision	recall	f1-score	support
1	1.00	1.00	1.00	11478
2	1.00	0.92	0.96	13
3	1.00	1.00	1.00	39
4	1.00	1.00	1.00	2155
5	1.00	1.00	1.00	809
6	1.00	1.00	1.00	4
7	1.00	1.00	1.00	2
accuracy			1.00	14500
macro avg	1.00	0.99	0.99	14500
weighted avg	1.00	1.00	1.00	14500

```
[[11478  0  0  0  0  0  0]
 [  0 12  0  1  0  0  0]
 [  0  0 39  0  0  0  0]
 [  0  0  0 2155  0  0  0]
 [  0  0  0  0 809  0  0]
 [  0  0  0  0  0  4  0]
 [  0  0  0  0  0  0  2]]
```

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                        max_depth=None, max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, presort='deprecated',
                        random_state=None, splitter='best')Accuracy: 1.00 (+/-
```

0.00)

Num Features: 9

Selected Features: [True True True True True True True True True]
 Feature Ranking: [1 1 1 1 1 1 1 1 1]

2.9.4 Random Forest

Let us evaluate the Random Forest algorithm on our second data set.

```
[14]: model(rf,X1_train,y1_train,X1_test,y1_test)
      crossval(rf,df_row_reindex_X,df_row_reindex_y)
      rfe(rf,df_row_reindex_X,df_row_reindex_y)
```

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=None, max_features='auto',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=100,
                        n_jobs=-1, oob_score=False, random_state=0, verbose=0,
                        warm_start=False)Accuracy: 99.97931034482758
```

	precision	recall	f1-score	support
1	1.00	1.00	1.00	11478
2	1.00	0.92	0.96	13
3	0.97	1.00	0.99	39
4	1.00	1.00	1.00	2155
5	1.00	1.00	1.00	809
6	1.00	0.75	0.86	4
7	1.00	0.50	0.67	2
accuracy			1.00	14500
macro avg	1.00	0.88	0.92	14500
weighted avg	1.00	1.00	1.00	14500

```
[[11478  0  0  0  0  0  0]
 [  0 12  0  1  0  0  0]
 [  0  0 39  0  0  0  0]
 [  0  0  0 2155  0  0  0]
 [  0  0  0  0 809  0  0]
 [  0  0  0  1  0  3  0]
 [  0  0  1  0  0  0  1]]
```

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=None, max_features='auto',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=100,
                        n_jobs=-1, oob_score=False, random_state=0, verbose=0,
                        warm_start=False)Accuracy: 1.00 (+/- 0.00)
```

Num Features: 9

Selected Features: [True True True True True True True True True]

Feature Ranking: [1 1 1 1 1 1 1 1 1]

2.9.5 Logistic Regression

Let us evaluate the Logistic Regression algorithm on our second data set.

```
[15]: model(logreg, X1_train,y1_train,X1_test,y1_test)
      crossval(logreg,df_row_reindex_X,df_row_reindex_y)
      rfe(logreg,df_row_reindex_X,df_row_reindex_y)
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='ovr', n_jobs=None, penalty='l2',
                    random_state=None, solver='liblinear', tol=0.0001, verbose=0,
                    warm_start=False)Accuracy: 93.10344827586206
```

	precision	recall	f1-score	support
1	0.93	0.99	0.96	11478
2	0.00	0.00	0.00	13
3	0.00	0.00	0.00	39
4	0.91	0.61	0.73	2155
5	1.00	1.00	1.00	809
6	0.00	0.00	0.00	4
7	0.00	0.00	0.00	2
accuracy			0.93	14500
macro avg	0.41	0.37	0.38	14500
weighted avg	0.93	0.93	0.92	14500

```
[[11372  0  0 104  0  0  2]
 [  8  0  0  5  0  0  0]
 [ 17  0  0 22  0  0  0]
 [ 835  0  0 1320  0  0  0]
 [  1  0  0  0 808  0  0]
 [  0  0  0  4  0  0  0]
 [  2  0  0  0  0  0  0]]
```

```
/home/kaygun/.local/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1268: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
/home/kaygun/.local/lib/python3.8/site-packages/sklearn/svm/_base.py:946:
ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
```

```
warnings.warn("Liblinear failed to converge, increase "
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='ovr', n_jobs=None, penalty='l2',
                    random_state=None, solver='liblinear', tol=0.0001, verbose=0,
                    warm_start=False)Accuracy: 0.93 (+/- 0.01)
```

Num Features: 9

Selected Features: [True True True True True True True True True]

Feature Ranking: [1 1 1 1 1 1 1 1 1]

2.9.6 Support Vector Machines

Let us evaluate the Support Vector Machines algorithm on our second data set. (My machine was unable to perform this operation.)

```
[ ]: model(svm,X1_train,y1_train,X1_test,y1_test)
      crossval(svm,df_row_reindex_X,df_row_reindex_y)
```

2.9.7 XGBoost

Let us evaluate the XGBoost algorithm on our second data set.

```
[16]: model(xgb,X1_train,y1_train,X1_test,y1_test)
      crossval(xgb,df_row_reindex_X,df_row_reindex_y)
      rfe(xgb,df_row_reindex_X,df_row_reindex_y)
```

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
               colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
               importance_type='gain', interaction_constraints='',
               learning_rate=0.300000012, max_delta_step=0, max_depth=6,
               min_child_weight=1, missing=nan, monotone_constraints='()',
               n_estimators=100, n_jobs=0, num_parallel_tree=1,
               objective='multi:softprob', random_state=0, reg_alpha=0,
               reg_lambda=1, scale_pos_weight=None, subsample=1,
               tree_method='exact', validate_parameters=1,
               verbosity=None)Accuracy: 99.99310344827586
```

	precision	recall	f1-score	support
1	1.00	1.00	1.00	11478
2	1.00	0.92	0.96	13
3	1.00	1.00	1.00	39
4	1.00	1.00	1.00	2155
5	1.00	1.00	1.00	809
6	1.00	1.00	1.00	4
7	1.00	1.00	1.00	2
accuracy			1.00	14500
macro avg	1.00	0.99	0.99	14500
weighted avg	1.00	1.00	1.00	14500

```
[[11478    0    0    0    0    0    0]
 [    0   12    0    1    0    0    0]
 [    0    0   39    0    0    0    0]
 [    0    0    0  2155    0    0    0]
 [    0    0    0    0   809    0    0]
 [    0    0    0    0    0    4    0]
 [    0    0    0    0    0    0    2]]
```

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
               colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
               importance_type='gain', interaction_constraints='',
               learning_rate=0.300000012, max_delta_step=0, max_depth=6,
               min_child_weight=1, missing=nan, monotone_constraints='()',
               n_estimators=100, n_jobs=0, num_parallel_tree=1,
               objective='multi:softprob', random_state=0, reg_alpha=0,
               reg_lambda=1, scale_pos_weight=None, subsample=1,
               tree_method='exact', validate_parameters=1,
               verbosity=None)Accuracy: 1.00 (+/- 0.00)
Num Features: 9
Selected Features: [ True  True  True  True  True  True  True  True  True]
Feature Ranking: [1 1 1 1 1 1 1 1 1]
```

2.9.8 Gaussian Naive Bayes

Let us evaluate the Gaussian Naive Bayes algorithm on our second data set.

```
[17]: model(gnb,X1_train,y1_train,X1_test,y1_test)
      crossval(gnb,df_row_reindex_X,df_row_reindex_y)
```

```
GaussianNB(priors=None, var_smoothing=1e-09)Accuracy: 82.6551724137931
```

	precision	recall	f1-score	support
1	0.95	0.88	0.92	11478
2	0.01	0.92	0.02	13
3	0.11	0.59	0.19	39
4	0.89	0.54	0.67	2155
5	0.99	0.82	0.90	809
6	0.40	1.00	0.57	4
7	0.00	1.00	0.01	2
accuracy			0.83	14500
macro avg	0.48	0.82	0.47	14500
weighted avg	0.94	0.83	0.88	14500

```
[[10116  463  185  143    4    5  562]
 [    1   12    0    0    0    0    0]
 [    7    1   23    0    1    0    7]
 [  502  491    0  1162    0    0    0]]
```



```
[ 0 142 0 0 666 1 0]
[ 0 0 0 0 0 4 0]
[ 0 0 0 0 0 0 2]]
```

GaussianNB(priors=None, var_smoothing=1e-09)Accuracy: 0.81 (+/- 0.01)

2.10 Experiments on Our Third Data Set

```
[20]: db = pd.read_csv("data-hastalikh.csv", sep = ';' )
db.head()
```

```
[20]:      ATES  BULANTI  BEL-AGRI  SUREKLI-WC  IDRAR-SIRASINDA-AGRI  \
0    355         0         1         0         0
1    359         0         0         1         1
2    359         0         1         0         0
3    360         0         0         1         1
4    360         0         1         0         0

      URETRADA-YANMA-SISME-KASINTI  MESANE-ILTIHABI  BOBREK-ILTIHABI
0                                0                0                0
1                                1                1                0
2                                0                0                0
3                                1                1                0
4                                0                0                0
```

2.10.1 Feature selection and splitting the train and test data sets

Let us denote the features we are going to use with X and response variable with y . With the drop function, we get the columns other than the columns ‘MESANE-ILTIHABI’ and ‘BOBREK-ILTIHABI’. For the response variable, we just take the column named ‘MESANE-ILTIHABI’. Then we determine the size of our training and test sets. For this dataset, 80% of all data is used for training and the remaining 20% is used for testing using the function “train_test_split”.

```
[50]: X = db.drop(["MESANE-ILTIHABI", "BOBREK-ILTIHABI"], axis=1)
y = db['MESANE-ILTIHABI']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2 )
```

```
[51]: A = db.drop(["MESANE-ILTIHABI", "BOBREK-ILTIHABI"], axis=1)
b = db['BOBREK-ILTIHABI']
A_train, A_test, b_train, b_test = train_test_split(A, b, test_size=0.2 )
```

2.10.2 Create array for loop

We create two arrays for our third dataset. We will use these arrays for the functions we will write below. The difference between the two arrays is that the rfe function does not work in SVM and GNB algorithms.

```
[ ]: models = [dt, rf, logreg, svm, xgb, gnb]
```

2.10.3 Evaluate Models

Let us evaluate the all model on our third data set.

```
[85]: for i in models:
      model(i, X_train, X_test, y_train, y_test)
      model(i, A_train, A_test, b_train, b_test)
      crossval(i, X, y, 5)
      crossval(i, X, y, 10)
      crossval(i, A, b, 5)
      crossval(i, A, b, 10)
```

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                       max_depth=None, max_features=None, max_leaf_nodes=None,
                       min_impurity_decrease=0.0, min_impurity_split=None,
                       min_samples_leaf=1, min_samples_split=2,
                       min_weight_fraction_leaf=0.0, presort='deprecated',
                       random_state=None, splitter='best')Accuracy: 100.0
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	14
1	1.00	1.00	1.00	10
accuracy			1.00	24
macro avg	1.00	1.00	1.00	24
weighted avg	1.00	1.00	1.00	24

```
[[14  0]
 [ 0 10]]
```

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                       max_depth=None, max_features=None, max_leaf_nodes=None,
                       min_impurity_decrease=0.0, min_impurity_split=None,
                       min_samples_leaf=1, min_samples_split=2,
                       min_weight_fraction_leaf=0.0, presort='deprecated',
                       random_state=None, splitter='best')Accuracy: 100.0
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	15
1	1.00	1.00	1.00	9
accuracy			1.00	24
macro avg	1.00	1.00	1.00	24
weighted avg	1.00	1.00	1.00	24

```
[[15  0]
 [ 0  9]]
```

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                       max_depth=None, max_features=None, max_leaf_nodes=None,
```

```

min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort='deprecated',
random_state=None, splitter='best')Accuracy: 0.97 (+/-
0.13)
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
max_depth=None, max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort='deprecated',
random_state=None, splitter='best')Accuracy: 1.00 (+/-
0.00)
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
max_depth=None, max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort='deprecated',
random_state=None, splitter='best')Accuracy: 0.86 (+/-
0.36)
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
max_depth=None, max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort='deprecated',
random_state=None, splitter='best')Accuracy: 0.97 (+/-
0.20)
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
criterion='gini', max_depth=None, max_features='auto',
max_leaf_nodes=None, max_samples=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=100,
n_jobs=-1, oob_score=False, random_state=0, verbose=0,
warm_start=False)Accuracy: 100.0
precision    recall  f1-score   support

0           1.00      1.00      1.00        14
1           1.00      1.00      1.00        10

accuracy               1.00        24
macro avg           1.00      1.00      1.00        24
weighted avg       1.00      1.00      1.00        24

[[14  0]
 [ 0 10]]
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
criterion='gini', max_depth=None, max_features='auto',
max_leaf_nodes=None, max_samples=None,

```

```

min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=100,
n_jobs=-1, oob_score=False, random_state=0, verbose=0,
warm_start=False)Accuracy: 100.0
precision    recall  f1-score   support

0           1.00      1.00      1.00        15
1           1.00      1.00      1.00         9

accuracy                1.00        24
macro avg              1.00      1.00      1.00        24
weighted avg          1.00      1.00      1.00        24

[[15  0]
 [ 0  9]]
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=None, max_features='auto',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=100,
                        n_jobs=-1, oob_score=False, random_state=0, verbose=0,
                        warm_start=False)Accuracy: 1.00 (+/- 0.00)
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=None, max_features='auto',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=100,
                        n_jobs=-1, oob_score=False, random_state=0, verbose=0,
                        warm_start=False)Accuracy: 1.00 (+/- 0.00)
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=None, max_features='auto',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=100,
                        n_jobs=-1, oob_score=False, random_state=0, verbose=0,
                        warm_start=False)Accuracy: 1.00 (+/- 0.00)
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=None, max_features='auto',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=100,
                        n_jobs=-1, oob_score=False, random_state=0, verbose=0,
                        warm_start=False)Accuracy: 0.97 (+/- 0.20)

```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='ovr', n_jobs=None, penalty='l2',
                    random_state=None, solver='liblinear', tol=0.0001, verbose=0,
                    warm_start=False)Accuracy: 100.0
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	14
1	1.00	1.00	1.00	10
accuracy			1.00	24
macro avg	1.00	1.00	1.00	24
weighted avg	1.00	1.00	1.00	24

```
[[14  0]
 [ 0 10]]
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='ovr', n_jobs=None, penalty='l2',
                    random_state=None, solver='liblinear', tol=0.0001, verbose=0,
                    warm_start=False)Accuracy: 100.0
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	15
1	1.00	1.00	1.00	9
accuracy			1.00	24
macro avg	1.00	1.00	1.00	24
weighted avg	1.00	1.00	1.00	24

```
[[15  0]
 [ 0  9]]
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='ovr', n_jobs=None, penalty='l2',
                    random_state=None, solver='liblinear', tol=0.0001, verbose=0,
                    warm_start=False)Accuracy: 1.00 (+/- 0.00)
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='ovr', n_jobs=None, penalty='l2',
                    random_state=None, solver='liblinear', tol=0.0001, verbose=0,
                    warm_start=False)Accuracy: 1.00 (+/- 0.00)
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='ovr', n_jobs=None, penalty='l2',
                    random_state=None, solver='liblinear', tol=0.0001, verbose=0,
                    warm_start=False)Accuracy: 1.00 (+/- 0.00)
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
```

```

        intercept_scaling=1, l1_ratio=None, max_iter=100,
        multi_class='ovr', n_jobs=None, penalty='l2',
        random_state=None, solver='liblinear', tol=0.0001, verbose=0,
        warm_start=False)Accuracy: 1.00 (+/- 0.00)
SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)Accuracy: 87.5
      precision    recall  f1-score   support

         0         1.00      0.79      0.88         14
         1         0.77      1.00      0.87         10

   accuracy
macro avg      0.88      0.89      0.87         24
weighted avg   0.90      0.88      0.88         24

[[11  3]
 [ 0 10]]
SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)Accuracy: 91.66666666666666
      precision    recall  f1-score   support

         0         1.00      0.87      0.93         15
         1         0.82      1.00      0.90          9

   accuracy
macro avg      0.91      0.93      0.91         24
weighted avg   0.93      0.92      0.92         24

[[13  2]
 [ 0  9]]
SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)Accuracy: 0.54 (+/- 0.17)
SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)Accuracy: 0.63 (+/- 0.33)
SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)Accuracy: 0.74 (+/- 0.33)
SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',

```

```

max_iter=-1, probability=False, random_state=None, shrinking=True,
tol=0.001, verbose=False)Accuracy: 0.84 (+/- 0.39)
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
               colsample_bynode=1, colsample_bytree=1, gamma=0,
               learning_rate=0.1, max_delta_step=0, max_depth=3,
               min_child_weight=1, missing=None, n_estimators=100, n_jobs=1,
               nthread=None, objective='binary:logistic', random_state=0,
               reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
               silent=None, subsample=1, verbosity=1)Accuracy: 100.0
precision    recall  f1-score   support

      0         1.00      1.00      1.00        14
      1         1.00      1.00      1.00        10

   accuracy                1.00        24
  macro avg         1.00      1.00      1.00        24
weighted avg         1.00      1.00      1.00        24

[[14  0]
 [ 0 10]]
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
               colsample_bynode=1, colsample_bytree=1, gamma=0,
               learning_rate=0.1, max_delta_step=0, max_depth=3,
               min_child_weight=1, missing=None, n_estimators=100, n_jobs=1,
               nthread=None, objective='binary:logistic', random_state=0,
               reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
               silent=None, subsample=1, verbosity=1)Accuracy: 100.0
precision    recall  f1-score   support

      0         1.00      1.00      1.00        15
      1         1.00      1.00      1.00         9

   accuracy                1.00        24
  macro avg         1.00      1.00      1.00        24
weighted avg         1.00      1.00      1.00        24

[[15  0]
 [ 0  9]]
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
               colsample_bynode=1, colsample_bytree=1, gamma=0,
               learning_rate=0.1, max_delta_step=0, max_depth=3,
               min_child_weight=1, missing=None, n_estimators=100, n_jobs=1,
               nthread=None, objective='binary:logistic', random_state=0,
               reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
               silent=None, subsample=1, verbosity=1)Accuracy: 0.97 (+/- 0.10)
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
               colsample_bynode=1, colsample_bytree=1, gamma=0,
               learning_rate=0.1, max_delta_step=0, max_depth=3,

```

```

min_child_weight=1, missing=None, n_estimators=100, n_jobs=1,
nthread=None, objective='binary:logistic', random_state=0,
reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
silent=None, subsample=1, verbosity=1)Accuracy: 0.97 (+/- 0.15)
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
colsample_bynode=1, colsample_bytree=1, gamma=0,
learning_rate=0.1, max_delta_step=0, max_depth=3,
min_child_weight=1, missing=None, n_estimators=100, n_jobs=1,
nthread=None, objective='binary:logistic', random_state=0,
reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
silent=None, subsample=1, verbosity=1)Accuracy: 0.86 (+/- 0.36)
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
colsample_bynode=1, colsample_bytree=1, gamma=0,
learning_rate=0.1, max_delta_step=0, max_depth=3,
min_child_weight=1, missing=None, n_estimators=100, n_jobs=1,
nthread=None, objective='binary:logistic', random_state=0,
reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
silent=None, subsample=1, verbosity=1)Accuracy: 0.97 (+/- 0.15)
GaussianNB(priors=None, var_smoothing=1e-09)Accuracy: 70.83333333333334
precision    recall  f1-score   support

      0         1.00      0.50      0.67         14
      1         0.59      1.00      0.74         10

   accuracy                   0.71         24
  macro avg          0.79      0.75      0.70         24
weighted avg          0.83      0.71      0.70         24

[[ 7  7]
 [ 0 10]]
GaussianNB(priors=None, var_smoothing=1e-09)Accuracy: 91.66666666666666
precision    recall  f1-score   support

      0         0.88      1.00      0.94         15
      1         1.00      0.78      0.88          9

   accuracy                   0.92         24
  macro avg          0.94      0.89      0.91         24
weighted avg          0.93      0.92      0.91         24

[[15  0]
 [ 2  7]]
GaussianNB(priors=None, var_smoothing=1e-09)Accuracy: 0.82 (+/- 0.19)
GaussianNB(priors=None, var_smoothing=1e-09)Accuracy: 0.82 (+/- 0.32)
GaussianNB(priors=None, var_smoothing=1e-09)Accuracy: 0.92 (+/- 0.33)
GaussianNB(priors=None, var_smoothing=1e-09)Accuracy: 0.95 (+/- 0.25)

```


2.10.4 Recursive Feature Elimination

Evaluate rfe on our third data set

```
[67]: model = [dt, rf, logreg, xgb]
      for i in model:
          print(i)
          rfe(i, X, y)
          rfe(i, A, b)
```

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                      max_depth=None, max_features=None, max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, presort='deprecated',
                      random_state=None, splitter='best')
```

Num Features: 2

Selected Features: [False False False True True False]

Feature Ranking: [2 5 4 1 1 3]

Num Features: 2

Selected Features: [True False True False False False]

Feature Ranking: [1 5 1 4 3 2]

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                      criterion='gini', max_depth=None, max_features='auto',
                      max_leaf_nodes=None, max_samples=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, n_estimators=100,
                      n_jobs=-1, oob_score=False, random_state=0, verbose=0,
                      warm_start=False)
```

Num Features: 2

Selected Features: [False False False True True False]

Feature Ranking: [3 5 2 1 1 4]

Num Features: 2

Selected Features: [True False True False False False]

Feature Ranking: [1 2 1 5 4 3]

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                   intercept_scaling=1, l1_ratio=None, max_iter=100,
                   multi_class='ovr', n_jobs=None, penalty='l2',
                   random_state=None, solver='liblinear', tol=0.0001, verbose=0,
                   warm_start=False)
```

Num Features: 2

Selected Features: [False False False True True False]

Feature Ranking: [5 3 2 1 1 4]

Num Features: 2

Selected Features: [False True True False False False]

Feature Ranking: [5 1 1 3 4 2]

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
```

```
colsample_bynode=1, colsample_bytree=1, gamma=0,  
learning_rate=0.1, max_delta_step=0, max_depth=3,  
min_child_weight=1, missing=None, n_estimators=100, n_jobs=1,  
nthread=None, objective='binary:logistic', random_state=0,  
reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,  
silent=None, subsample=1, verbosity=1)
```

Num Features: 2

Selected Features: [False False False True True False]

Feature Ranking: [2 3 4 1 1 5]

Num Features: 2

Selected Features: [True False True False False False]

Feature Ranking: [1 5 1 2 3 4]